

wrangle_act2

January 26, 2018

1 WeRateDogs Data Wrangling

The dataset that we have been wrangling consists of three sources:

the tweet archive of Twitter user @dog_rates He is also known as WeRateDogs. WeRateDogs is a Twitter account that rates people's dogs with a humorous comment about the dog. These ratings almost always have a denominator of 10. The numerators, though? Almost always greater than 10. 11/10, 12/10, 13/10, etc. Why? Because "they're good dogs Brent." WeRateDogs has over 4 million followers and has received international media coverage.

Additional Data via the Twitter API

Back to the basic-ness of Twitter archives: retweet count and favorite count are two of the notable column omissions. Fortunately, this additional data can be gathered by anyone from Twitter's API. Well, "anyone" who has access to data for the 3000 most recent tweets, at least. But you, because you have the WeRateDogs Twitter archive and specifically the tweet IDs within it, can gather this data for all 5000+. And guess what? You're going to query Twitter's API to gather this valuable data.

The current Tweet data of those archived tweets has been extracted from the life twitter API

-

1.1 Gather

```
In [2]: import requests
```

```
#deal with data
import numpy as np
import pandas as pd

#deal with datetime
import datetime as dt
import pytz

#deal with visualization
import seaborn as sns
%matplotlib inline
import matplotlib.pyplot as plt

#use padasql for SQL-query on dataframe
```

```
#http://blog.yhat.com/posts/pandasql-intro.html
from pandasql import sqldf
```

1.1.1 gather twitter-archive-enhanced.csv file

```
In [2]: #read locally stored file
df_twarchive = pd.read_csv('twitter-archive-enhanced.csv')
```

1.1.2 gather image-predictions.tsv file

```
In [3]: tgt_filename = 'image-predictions.tsv'
```

```
In [15]: #download file from the cloud
url = 'https://d17h27t6h515a5.cloudfront.net/topher/2017/August/599fd2ad_image-predictions.tsv'
r = requests.get(url)
with open(tgt_filename, 'wb') as file:
    file.write(r.content)
```

```
In [4]: #read dataframe from file
df_image = pd.read_csv(tgt_filename, sep = '\t')
df_image.tweet_id.count()
```

```
Out[4]: 2075
```

1.1.3 gather tweed_json.txt

Each tweet's retweet count and favorite ("like") count at minimum, and any additional data you find interesting. Using the **tweet IDs in the WeRateDogs Twitter archive**, query the Twitter API for each tweet's JSON data using Python's Tweepy library and store each tweet's entire set of JSON data in a file called tweet_json.txt file. Each tweet's JSON data should be written to its own line. Then read this .txt file line by line into a pandas DataFrame with (at minimum) - tweet ID - retweet count - favorite count

```
In [5]: #refer to twitter application https://apps.twitter.com/app/14714948/keys
#import authorization variables for security reasons
import access.my_secret as xs
import tweepy

auth = tweepy.OAuthHandler(xs.consumer_key, xs.consumer_secret)
auth.set_access_token(xs.access_token, xs.access_secret)
api = tweepy.API(auth)
```

Extract tweet in json format from twitter api: `def get_tweet(tweet_id): tweet = api.get_status(tweet_id, tweet_mode='extended', wait_on_rate_limit=True) return tweet._json`

```
In [20]: # for testing
# print(get_tweet(666082916733198337))
```

```

In [5]: #get the timer set up
        from timeit import default_timer as timer
        import datetime
        import json

        json_filename = 'tweet_json.txt'

In [22]: #get tweet
        #write json file line by line and log progress
        with open(json_filename, 'w') as file, open('log.txt', 'w') as log:
            for tweet_id in df_twarchive.tweet_id:
                result = ''
                start = timer()
                try:
                    content = get_tweet(tweet_id)
                    #add newline is helpful to read the file line by line later
                    file.write(json.dumps(content) + '\n')
                    result = 'ok'
                except tweepy.TweepError:
                    result = 'TweepError'
                end = timer()
                log.write('%s\t%s\t%s\t%s\n' % (str(datetime.datetime.now()), result, str(tweet_id), end - start))

In [6]: #count number of tweets
        i = 0
        with open(json_filename, 'r') as file:
            for row in file:
                i += 1
        print("row count:", i)

row count: 2345

```

1.2 Assess

After gathering each of the above pieces of data, assess them visually and programmatically for quality and tidiness issues. Detect and document at least **eight (8) quality issues and two (2) tidiness issues** in your wrangle_act.ipynb Jupyter Notebook. To meet specifications, the **issues that satisfy the Project Motivation must be assessed**.

Of the 5000+ tweets, I have filtered for tweets with ratings only (there are 2356).

1.2.1 assess twitter-archive-enhanced.csv file

quality issues

1. derive clear categories for source(iPhone App, vine.co, Twitter Web Client, TweetDeck)
2. make this source column categorical
3. deal with a few invalid ratings where rating_numerator is 0 or rating_denominator is 0

- 835246439529840640 rating has been changed to from 960/0 to 13/10 according to current twitter post
 - 835152434251116546 rating has been changed to from 0/10 to 11/10 according to current twitter post
 - 746906459439529985 no rating according to current twitter post -> filter out
4. take out ratings with a rating_denominator other than 10 as these are invalid ratings or rate multiple dogs at once.
 5. take out ratings with rating_numerator > 15 they seem to be invalid
 6. take out retweets You only want original ratings (no retweets) that have images. Though there are 5000+ tweets in the dataset, not all are dog ratings and some are retweets. (retweeted_status_id is NaN)
 7. convert wrong dog names to Null: "a", "the", "an", "n", "None"

tidyness issues

1. four columns for the dog_stage can be put into one categorial variable (doggo,floofer,pupper,puppo) and turn the default value "None" into a Null or "unknown"

```
In [7]: df_twarchive.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2356 non-null int64
in_reply_to_status_id    78 non-null float64
in_reply_to_user_id      78 non-null float64
timestamp               2356 non-null object
source                  2356 non-null object
text                    2356 non-null object
retweeted_status_id      181 non-null float64
retweeted_status_user_id 181 non-null float64
retweeted_status_timestamp 181 non-null object
expanded_urls            2297 non-null object
rating_numerator         2356 non-null int64
rating_denominator       2356 non-null int64
name                    2356 non-null object
doggo                   2356 non-null object
floofer                 2356 non-null object
pupper                  2356 non-null object
puppo                   2356 non-null object
dtypes: float64(4), int64(3), object(10)
memory usage: 313.0+ KB
```

```
In [8]: df_twarchive.sample(10)
```

```
Out[8]:
```

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	\
1473	693647888581312512	NaN	NaN	

2320	666437273139982337	NaN	NaN
1690	681339448655802368	NaN	NaN
1100	735648611367784448	NaN	NaN
28	886680336477933568	NaN	NaN
1326	706153300320784384	NaN	NaN
559	803276597545603072	NaN	NaN
2071	671115716440031232	NaN	NaN
802	772193107915964416	NaN	NaN
2125	670361874861563904	NaN	NaN

	timestamp \
1473	2016-01-31 04:11:58 +0000
2320	2015-11-17 02:06:42 +0000
1690	2015-12-28 05:02:37 +0000
1100	2016-05-26 01:47:51 +0000
28	2017-07-16 20:14:00 +0000
1326	2016-03-05 16:24:01 +0000
559	2016-11-28 16:37:19 +0000
2071	2015-11-29 23:57:10 +0000
802	2016-09-03 22:02:38 +0000
2125	2015-11-27 22:01:40 +0000

	source \
1473	<a href="http://twitter.com/download/iphone" r...
2320	<a href="http://twitter.com/download/iphone" r...
1690	<a href="http://twitter.com/download/iphone" r...
1100	<a href="http://twitter.com/download/iphone" r...
28	<a href="http://twitter.com/download/iphone" r...
1326	Vine -...
559	<a href="http://twitter.com/download/iphone" r...
2071	<a href="http://twitter.com/download/iphone" r...
802	<a href="http://twitter.com/download/iphone" r...
2125	<a href="http://twitter.com/download/iphone" r...

	text	retweeted_status_id \
1473	What kind of person sends in a pic without a d...	NaN
2320	Here we see a lone northeastern Cumberbatch. H...	NaN
1690	This is Asher. He's not wearing a seatbelt or ...	NaN
1100	*faints* 12/10 perfection in pupper form https...	NaN
28	This is Derek. He's late for a dog meeting. 13...	NaN
1326	This is Layla. She's giving you a standing ova...	NaN
559	This is Winston. His selfie game is legendary...	NaN
2071	Meet Phred. He isn't steering, looking at the ...	NaN
802	This is Willem. He's a Penn State pupper. Thin...	NaN
2125	This is a Rich Mahogany Seltzer named Cherokee...	NaN

	retweeted_status_user_id	retweeted_status_timestamp \
1473	NaN	NaN

2320	NaN	NaN
1690	NaN	NaN
1100	NaN	NaN
28	NaN	NaN
1326	NaN	NaN
559	NaN	NaN
2071	NaN	NaN
802	NaN	NaN
2125	NaN	NaN

	expanded_urls	rating_numerator	\
1473	https://twitter.com/dog_rates/status/693647888...	7	
2320	https://twitter.com/dog_rates/status/666437273...	7	
1690	https://twitter.com/dog_rates/status/681339448...	9	
1100	https://twitter.com/dog_rates/status/735648611...	12	
28	https://twitter.com/dog_rates/status/886680336...	13	
1326	https://vine.co/v/iXidJXBJ3P9	13	
559	https://twitter.com/dog_rates/status/803276597...	11	
2071	https://twitter.com/dog_rates/status/671115716...	6	
802	https://twitter.com/dog_rates/status/772193107...	12	
2125	https://twitter.com/dog_rates/status/670361874...	9	

	rating_denominator	name	doggo	floofer	pupper	puppo
1473	10	None	None	None	None	None
2320	10	None	None	None	None	None
1690	10	Asher	None	None	None	None
1100	10	None	None	None	pupper	None
28	10	Derek	None	None	None	None
1326	10	Layla	None	None	None	None
559	10	Winston	None	None	None	None
2071	10	Phred	None	None	None	None
802	10	Willem	None	None	pupper	None
2125	10	a	None	None	None	None

In [9]: df_twarchive.source.value_counts()

Out[9]: Twitter for iPhone
 Vine - Make a Scene
 Twitter Web Client
 TweetDeck
 Name: source, dtype: int64

In [10]: *#check the latest tweet*
 print(df_twarchive.timestamp.min())
 print(df_twarchive.timestamp.max())

2015-11-15 22:32:08 +0000
 2017-08-01 16:23:56 +0000

```
In [11]: df_twarchive.describe()
```

```
Out[11]:
```

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	\
count	2.356000e+03	7.800000e+01	7.800000e+01	
mean	7.427716e+17	7.455079e+17	2.014171e+16	
std	6.856705e+16	7.582492e+16	1.252797e+17	
min	6.660209e+17	6.658147e+17	1.185634e+07	
25%	6.783989e+17	6.757419e+17	3.086374e+08	
50%	7.196279e+17	7.038708e+17	4.196984e+09	
75%	7.993373e+17	8.257804e+17	4.196984e+09	
max	8.924206e+17	8.862664e+17	8.405479e+17	

	retweeted_status_id	retweeted_status_user_id	rating_numerator	\
count	1.810000e+02	1.810000e+02	2356.000000	
mean	7.720400e+17	1.241698e+16	13.126486	
std	6.236928e+16	9.599254e+16	45.876648	
min	6.661041e+17	7.832140e+05	0.000000	
25%	7.186315e+17	4.196984e+09	10.000000	
50%	7.804657e+17	4.196984e+09	11.000000	
75%	8.203146e+17	4.196984e+09	12.000000	
max	8.874740e+17	7.874618e+17	1776.000000	

	rating_denominator
count	2356.000000
mean	10.455433
std	6.745237
min	0.000000
25%	10.000000
50%	10.000000
75%	10.000000
max	170.000000

```
In [12]: df_twarchive.query('rating_numerator == 0 | rating_denominator == 0')
#835246439529840640 rating has been changed to from 960/0 to 13/10 according to current
#835152434251116546 rating has been changed to from 0/10 to 11/10 according to current
#746906459439529985 no rating according to current twitter post -> filter out
```

```
Out[12]:
```

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	\
313	835246439529840640	8.352460e+17	2.625958e+07	
315	835152434251116546	NaN	NaN	
1016	746906459439529985	7.468859e+17	4.196984e+09	

	timestamp	\
313	2017-02-24 21:54:03 +0000	
315	2017-02-24 15:40:31 +0000	
1016	2016-06-26 03:22:31 +0000	

```
source \
```

```

313 <a href="http://twitter.com/download/iphone" r...
315 <a href="http://twitter.com/download/iphone" r...
1016 <a href="http://twitter.com/download/iphone" r...

                                text  retweeted_status_id  \
313 @jonnyusun @Lin_Manuel ok jomny I know you're e...      NaN
315 When you're so blinded by your systematic plag...      NaN
1016 PUPDATE: can't see any. Even if I could, I cou...      NaN

    retweeted_status_user_id retweeted_status_timestamp  \
313                        NaN                        NaN
315                        NaN                        NaN
1016                       NaN                       NaN

                                expanded_urls  rating_numerator  \
313                                           NaN                960
315 https://twitter.com/dog_rates/status/835152434...           0
1016 https://twitter.com/dog_rates/status/746906459...           0

    rating_denominator  name  doggo  floofer  pupper  puppo
313                   0  None   None    None   None   None
315                  10  None   None    None   None   None
1016                 10  None   None    None   None   None

```

```
In [13]: df_twarchive[['tweet_id', 'rating_numerator', 'rating_denominator']].query('rating_denomi
```

```

Out[13]:
   tweet_id  rating_numerator  rating_denominator
313  835246439529840640         960              0
342  832088576586297345          11             15
433  820690176645140481          84             70
516  810984652412424192          24              7
784  775096608509886464           9             11
902  758467244762497024        165            150
1068  740373189193256964           9             11
1120  731156023742988288        204            170
1165  722974582966214656           4             20
1202  716439118184652801          50             50
1228  713900603437621249          99             90
1254  710658690886586372          80             80
1274  709198395643068416          45             50
1351  704054845121142784          60             50
1433  697463031882764288          44             40
1598  686035780142297088           4             20
1634  684225744407494656        143            130
1635  684222868335505415        121            110
1662  682962037429899265           7             11
1663  682808988178739200          20             16
1779  677716515794329600        144            120

```


1843	675853064436391936	88	80
2335	666287406224695296	1	2

```
In [14]: df = df_twarchive[['tweet_id', 'rating_numerator', 'rating_denominator']].query('rating_d
df.rating_numerator.value_counts()
df.query('rating_numerator > 15')
```

```
Out[14]:
```

	tweet_id	rating_numerator	rating_denominator
55	881633300179243008	17	10
188	855862651834028034	420	10
189	855860136149123072	666	10
290	838150277551247360	182	10
340	832215909146226688	75	10
695	786709082849828864	75	10
763	778027034220126208	27	10
979	749981277374128128	1776	10
1712	680494726643068929	26	10
2074	670842764863651840	420	10

```
In [15]: #check if categories are clean
categories = ['doggo', 'floofer', 'pupper', 'puppo']
for category in categories:
    print(df_twarchive[category].value_counts())
    print()
```

```
None      2259
doggo      97
Name: doggo, dtype: int64
```

```
None      2346
floofer    10
Name: floofer, dtype: int64
```

```
None      2099
pupper     257
Name: pupper, dtype: int64
```

```
None      2326
puppo      30
Name: puppo, dtype: int64
```

```
In [20]: df_twarchive.name.value_counts().head(20)
#convert to Null: a, the, an
```

```
Out[20]:
```

None	745
a	55
Charlie	12

Cooper	11
Oliver	11
Lucy	11
Tucker	10
Lola	10
Penny	10
Winston	9
Bo	9
the	8
Sadie	8
Bailey	7
an	7
Daisy	7
Toby	7
Buddy	7
Oscar	6
Koda	6

Name: name, dtype: int64

The tweet image predictions, i.e., what breed of dog (or other object, animal, etc.) is present in each tweet according to a neural network.

Example for the data set: - tweet_id is the last part of the tweet URL after “status/” https://twitter.com/dog_rates/status/889531135344209921 - p1 is the algorithm’s #1 prediction for the image in the tweet golden retriever - p1_conf is how confident the algorithm is in its #1 prediction 95% - p1_dog is whether or not the #1 prediction is a breed of dog TRUE - p2 is the algorithm’s second most likely prediction Labrador retriever - p2_conf is how confident the algorithm is in its #2 prediction 1% - p2_dog is whether or not the #2 prediction is a breed of dog TRUE - etc.

quality issues

- we have 66 duplicates and have to make sure they don’t exists in the final data set

tidiness issues

2. get the best true dog prediction in one column
3. make the dog prediction categorial

In [23]: df_image.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
tweet_id      2075 non-null int64
jpg_url       2075 non-null object
img_num       2075 non-null int64
p1            2075 non-null object
p1_conf       2075 non-null float64
p1_dog        2075 non-null bool
```

```

p2          2075 non-null object
p2_conf     2075 non-null float64
p2_dog      2075 non-null bool
p3          2075 non-null object
p3_conf     2075 non-null float64
p3_dog      2075 non-null bool
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 152.1+ KB

```

```
In [24]: df_image.sample(5)
```

```

Out[24]:
      tweet_id      jpg_url \
250  670691627984359425  https://pbs.twimg.com/media/CU7GehOUYAA9nn-.jpg
1128 727685679342333952  https://pbs.twimg.com/media/ChlCQg-VIAQ_8g4.jpg
1335 758041019896193024  https://pbs.twimg.com/media/CoUaSKEXYAAySAl.jpg
280  671115716440031232  https://pbs.twimg.com/media/CVBILUgVAAA1ZUr.jpg
858  696886256886657024  https://pbs.twimg.com/media/CavWWdFWAAArflW.jpg

      img_num      p1      p1_conf      p1_dog      p2 \
250         1  Shetland_sheepdog  0.071124      True      home_theater
1128         1    Border_collie  0.462408      True          collie
1335         1      bookshop  0.794272     False      Cardigan
280         1      malinois  0.406341      True          kelpie
858         1      kuvasz  0.383941      True  golden_retriever

      p2_conf      p2_dog      p3      p3_conf      p3_dog
250  0.068398     False  American_Staffordshire_terrier  0.066964      True
1128  0.214556      True                      Eskimo_dog  0.035604      True
1335  0.051265      True          Bernese_mountain_dog  0.026596      True
280  0.143366      True                      dingo  0.129802     False
858  0.289085      True                      dingo  0.056548     False

```

```
In [25]: df_image.describe()
```

```

Out[25]:
      tweet_id      img_num      p1_conf      p2_conf      p3_conf
count  2.075000e+03  2075.000000  2075.000000  2.075000e+03  2.075000e+03
mean    7.384514e+17    1.203855    0.594548    1.345886e-01  6.032417e-02
std     6.785203e+16    0.561875    0.271174    1.006657e-01  5.090593e-02
min     6.660209e+17    1.000000    0.044333    1.011300e-08  1.740170e-10
25%     6.764835e+17    1.000000    0.364412    5.388625e-02  1.622240e-02
50%     7.119988e+17    1.000000    0.588230    1.181810e-01  4.944380e-02
75%     7.932034e+17    1.000000    0.843855    1.955655e-01  9.180755e-02
max     8.924206e+17    4.000000    1.000000    4.880140e-01  2.734190e-01

```

```
In [50]: df_image[df_image.duplicated(subset=['jpg_url'], keep='first')].head(5)
```

NameError

Traceback (most recent call last)

```
<ipython-input-50-6f41e28b1ff4> in <module>()
----> 1 df_image[df_image.duplicated(subset=['jpg_url'], keep='first')].head(5)
```

NameError: name 'df_image' is not defined

1.2.2 assess tweet_json.txt

To assess the JSON object I refer to - <https://developer.twitter.com/en/docs/tweets/data-dictionary/overview/tweet-object> or - http://support.gnip.com/sources/twitter/data_format.html with more detail on - <http://support.gnip.com/doing-more-with-140.html> tweet ID, retweet count, and favorite count

```
In [27]: #read json file line by line
df_list=[]
with open(json_filename, 'r') as file:
    for row in file:
        data = json.loads(row)
        df_list.append({"tweet_id":data["id"],
                        "retweet_count":data["retweet_count"],
                        "favorite_count":data["favorite_count"]
                        })

#df_tweet = pd.DataFrame.from_dict(df_list)
df_tweet = pd.DataFrame(df_list, columns=["tweet_id", "retweet_count", "favorite_count"])
```

```
In [28]: df_tweet.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2345 entries, 0 to 2344
Data columns (total 3 columns):
tweet_id      2345 non-null int64
retweet_count 2345 non-null int64
favorite_count 2345 non-null int64
dtypes: int64(3)
memory usage: 55.0 KB
```

```
In [29]: df_tweet.describe()
```

```
Out[29]:
```

	tweet_id	retweet_count	favorite_count
count	2.345000e+03	2345.000000	2345.000000
mean	7.423760e+17	3078.081876	8114.238380
std	6.836820e+16	5105.627191	12189.503552
min	6.660209e+17	0.000000	0.000000
25%	6.783802e+17	614.000000	1410.000000

50%	7.189719e+17	1437.000000	3574.000000
75%	7.987057e+17	3585.000000	10068.000000
max	8.924206e+17	78450.000000	142572.000000

```
In [30]: df_tweet[df_tweet.retweet_count>70000]
```

```
Out[30]:
```

	tweet_id	retweet_count	favorite_count
1029	744234799360020481	78450	129628

```
In [32]: df_twarchive[df_twarchive.tweet_id == 744234799360020481]
```

```
Out[32]:
```

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	\
1039	744234799360020481		NaN	NaN

	timestamp	\
1039	2016-06-18 18:26:18 +0000	

	source	\
1039	<a href="http://twitter.com/download/iphone" r...	

	text	retweeted_status_id	\
1039	Here's a doggo realizing you can stand in a po...		NaN

	retweeted_status_user_id	retweeted_status_timestamp	\
1039	NaN		NaN

	expanded_urls	rating_numerator	\
1039	https://twitter.com/dog_rates/status/744234799...		13

	rating_denominator	name	doggo	floofer	pupper	puppo
1039	10	None	doggo	None	None	None

1.3 Assess new file twitter_archive_master.csv

tidyness issues

- merge all three dataset into one
- drop columns in the merged data set which are not required

1.4 Clean

Clean each of the issues you documented while assessing. Perform this cleaning in wrangle_act.ipynb as well. The result should be a high quality and tidy master pandas DataFrame (or DataFrames, if appropriate). Again, **the issues that satisfy the Project Motivation must be cleaned.**

Clean `df_twarchive = pd.read_csv('twitter-archive-enhanced.csv')`

```
In [34]: #preserve the raw data
dfc_twarchive = df_twarchive.copy()
```

Define

- generate mapping dictionary to map to a short description with clear categories iPhone App, vine.co, Twitter Web Client, TweetDeck
- create function to change source
- call function
- change datatype to categorial

Code

```
In [35]: dfc_twarchive.source.value_counts()
```

```
Out[35]: <a href="http://twitter.com/download/iphone" rel="nofollow">Twitter for iPhone</a>
<a href="http://vine.co" rel="nofollow">Vine - Make a Scene</a>
<a href="http://twitter.com" rel="nofollow">Twitter Web Client</a>
<a href="https://about.twitter.com/products/tweetdeck" rel="nofollow">TweetDeck</a>
Name: source, dtype: int64
```

```
In [36]: ref_source = {'<a href="http://twitter.com/download/iphone" rel="nofollow">Twitter for
                        '<a href="http://vine.co" rel="nofollow">Vine - Make a Scene</a>': 'vine.
                        '<a href="http://twitter.com" rel="nofollow">Twitter Web Client</a>': 'Tw
                        '<a href="https://about.twitter.com/products/tweetdeck" rel="nofollow">Tw
```

```
def clean_source(tweet):
    if tweet['source'] in ref_source.keys():
        #return clean description
        return ref_source[tweet['source']]
    else:
        return dfc_twarchive.source
```

```
dfc_twarchive.source = dfc_twarchive.apply(clean_source, axis=1)
```

```
dfc_twarchive.source = dfc_twarchive.source.astype('category')
```

Test

```
In [37]: dfc_twarchive.source.value_counts()
```

```
Out[37]: iPhone App          2221
vine.co                      91
Twitter Web Client          33
TweetDeck                   11
Name: source, dtype: int64
```

```
In [38]: dfc_twarchive.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
```

```

tweet_id                2356 non-null int64
in_reply_to_status_id    78 non-null float64
in_reply_to_user_id      78 non-null float64
timestamp                2356 non-null object
source                   2356 non-null category
text                     2356 non-null object
retweeted_status_id      181 non-null float64
retweeted_status_user_id 181 non-null float64
retweeted_status_timestamp 181 non-null object
expanded_urls            2297 non-null object
rating_numerator          2356 non-null int64
rating_denominator        2356 non-null int64
name                     2356 non-null object
doggo                    2356 non-null object
floofer                  2356 non-null object
pupper                   2356 non-null object
puppo                    2356 non-null object
dtypes: category(1), float64(4), int64(3), object(9)
memory usage: 297.1+ KB

```

Define

- deal with a few invalid ratings where rating_numerator is 0 or rating_denominator is 0
 - 835246439529840640 rating has been changed to from 960/0 to 13/10 according to current twitter post
 - 835152434251116546 rating has been changed to from 0/10 to 11/10 according to current twitter post
 - 746906459439529985 no rating according to current twitter post -> filter out

Code

```

In [39]: #update df cell
         #see https://stackoverflow.com/questions/12307099/modifying-a-subset-of-rows-in-a-panda
dfc_twarchive.loc[dfc_twarchive.tweet_id == 835246439529840640, 'rating_numerator'] = 13
dfc_twarchive.loc[dfc_twarchive.tweet_id == 835246439529840640, 'rating_denominator'] = 10

dfc_twarchive.loc[dfc_twarchive.tweet_id == 835152434251116546, 'rating_numerator'] = 11
dfc_twarchive.loc[dfc_twarchive.tweet_id == 835152434251116546, 'rating_denominator'] = 10

dfc_twarchive = dfc_twarchive[dfc_twarchive.tweet_id != 746906459439529985]

```

Test

```

In [40]: dfc_twarchive.loc[dfc_twarchive.tweet_id.isin([835246439529840640, 835152434251116546,

```

```

Out[40]:
   tweet_id  in_reply_to_status_id  in_reply_to_user_id \
313  835246439529840640           8.352460e+17      26259576.0
315  835152434251116546                NaN                NaN

```

	timestamp	source	\
313	2017-02-24 21:54:03 +0000	iPhone App	
315	2017-02-24 15:40:31 +0000	iPhone App	

	text	retweeted_status_id	\
313	@jonnysun @Lin_Manuel ok jomny I know you're e...	NaN	
315	When you're so blinded by your systematic plag...	NaN	

	retweeted_status_user_id	retweeted_status_timestamp	\
313	NaN	NaN	
315	NaN	NaN	

	expanded_urls	rating_numerator	\
313	NaN	13	
315	https://twitter.com/dog_rates/status/835152434...	11	

	rating_denominator	name	doggo	floofer	pupper	puppo
313	10	None	None	None	None	None
315	10	None	None	None	None	None

Define

- take out ratings with a rating_denominator other than 10 as these are invalid ratings or rate multiple dogs at once.
- take out ratings with rating_numerator > 15 they seem to be invalid
- take out retweets You only want original ratings (no retweets) that have images. Though there are 5000+ tweets in the dataset, not all are dog ratings and some are retweets. So **keep** all tweets with retweeted_status_id is NaN

Code

```
In [41]: #delete retweets
dfc_twarchive = dfc_twarchive[dfc_twarchive.retweeted_status_id.isnull()]

In [42]: dfc_twarchive = dfc_twarchive[dfc_twarchive.rating_denominator == 10]
dfc_twarchive = dfc_twarchive[dfc_twarchive.rating_numerator <= 15]
```

Test

```
In [43]: dfc_twarchive.describe()
```

```
Out[43]:
```

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	\
count	2.144000e+03	6.900000e+01	6.900000e+01	
mean	7.370728e+17	7.403786e+17	2.276889e+16	
std	6.754280e+16	7.383460e+16	1.330848e+17	
min	6.660209e+17	6.658147e+17	1.185634e+07	
25%	6.766166e+17	6.753494e+17	1.198989e+09	
50%	7.095381e+17	7.030419e+17	4.196984e+09	

75%	7.895480e+17	8.131273e+17	4.196984e+09
max	8.924206e+17	8.862664e+17	8.405479e+17

	retweeted_status_id	retweeted_status_user_id	rating_numerator \
count	0.0	0.0	2144.000000
mean	NaN	NaN	10.630131
std	NaN	NaN	2.171990
min	NaN	NaN	1.000000
25%	NaN	NaN	10.000000
50%	NaN	NaN	11.000000
75%	NaN	NaN	12.000000
max	NaN	NaN	15.000000

	rating_denominator
count	2144.0
mean	10.0
std	0.0
min	10.0
25%	10.0
50%	10.0
75%	10.0
max	10.0

Define

- convert wrong dog names to Null: “a”, “the”, “an”, “n”

Code

```
In [44]: # Deal with value=None
# see https://stackoverflow.com/questions/17097236/how-to-replace-values-with-none-in-p
dfc_twarchive.name = dfc_twarchive.name.replace(to_replace=['^(a|an|the|n|None)$'], val
```

Test

```
In [45]: dfc_twarchive.loc[dfc_twarchive.name.isin(['a', 'an', 'the', 'n', 'None'])].name.count()
```

```
Out[45]: 0
```

Define

tidyness issues

- four columns for the dog_stage can be put into one categorial variable (doggo,floofer,pupper,puppo) and turn the default value “None” into a Null or “unknown”

Code

```
In [46]: dfc_twarchive.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2144 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2144 non-null int64
in_reply_to_status_id   69 non-null float64
in_reply_to_user_id     69 non-null float64
timestamp               2144 non-null object
source                  2144 non-null category
text                    2144 non-null object
retweeted_status_id      0 non-null float64
retweeted_status_user_id 0 non-null float64
retweeted_status_timestamp 0 non-null object
expanded_urls           2093 non-null object
rating_numerator         2144 non-null int64
rating_denominator       2144 non-null int64
name                    1419 non-null object
doggo                   2144 non-null object
floofer                 2144 non-null object
pupper                  2144 non-null object
puppo                   2144 non-null object
dtypes: category(1), float64(4), int64(3), object(9)
memory usage: 287.0+ KB
```

```
In [47]: #concat the stage columns doggo,floofer,pupper,puppo while replacing the None wordw with
dfc_twarchive['dog_stage'] = dfc_twarchive.doggo.replace('None','') + dfc_twarchive.floofer
dfc_twarchive['dog_stage'] = dfc_twarchive['dog_stage'].replace(to_replace='', value=)
```

```
In [48]: dfc_twarchive['dog_stage'].value_counts()
```

```
Out[48]: pupper          223
doggo                   75
puppo                   24
doggopupper            10
floofer                 9
doggopuppo              1
doggofloofer            1
Name: dog_stage, dtype: int64
```

```
In [49]: # get rid of dog stages which are not defined
# https://stackoverflow.com/questions/23330654/update-a-dataframe-in-pandas-while-iterating
for row in dfc_twarchive.itertuples():
    if dfc_twarchive.loc[row.Index, 'dog_stage'] in ['doggo','floofer','pupper','puppo']:
        next
    else:
        dfc_twarchive.loc[row.Index, 'dog_stage'] = None
dfc_twarchive['dog_stage'].value_counts()
```

```
Out[49]: pupper      223
         doggo       75
         puppo       24
         floofer      9
         Name: dog_stage, dtype: int64
```


```
In [50]: dfc_twarchive['dog_stage'] = dfc_twarchive['dog_stage'].astype('category')
```

Test

```
In [51]: dfc_twarchive.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2144 entries, 0 to 2355
Data columns (total 18 columns):
tweet_id                2144 non-null int64
in_reply_to_status_id   69 non-null float64
in_reply_to_user_id     69 non-null float64
timestamp               2144 non-null object
source                  2144 non-null category
text                   2144 non-null object
retweeted_status_id     0 non-null float64
retweeted_status_user_id 0 non-null float64
retweeted_status_timestamp 0 non-null object
expanded_urls           2093 non-null object
rating_numerator        2144 non-null int64
rating_denominator      2144 non-null int64
name                   1419 non-null object
doggo                  2144 non-null object
floofer                2144 non-null object
pupper                2144 non-null object
puppo                  2144 non-null object
dog_stage              331 non-null category
dtypes: category(2), float64(4), int64(3), object(9)
memory usage: 369.3+ KB
```

```
In [52]: sns.countplot(y="dog_stage", data=dfc_twarchive, palette="Greens_d");
```



Define drop all irrelevant columns

Code

```
In [53]: dfc_twarchive.drop(['in_reply_to_status_id',
                             'in_reply_to_user_id',
                             'retweeted_status_id',
                             'retweeted_status_user_id',
                             'retweeted_status_timestamp',
                             'doggo',
                             'floofer',
                             'pupper',
                             'puppo'], axis=1, inplace=True)
```

Test

```
In [54]: dfc_twarchive.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2144 entries, 0 to 2355
Data columns (total 9 columns):
tweet_id          2144 non-null int64
timestamp         2144 non-null object
source            2144 non-null category
```

```

text                2144 non-null object
expanded_urls       2093 non-null object
rating_numerator    2144 non-null int64
rating_denominator  2144 non-null int64
name                1419 non-null object
dog_stage           331 non-null category
dtypes: category(2), int64(3), object(4)
memory usage: 218.6+ KB

```

1.4.1 clean image-predictions.tsv file

```
In [55]: dfc_image = df_image.copy()
```

Define

- Go through each line and look for true dog predictions from column 1 to column 3.
- Stop at the first occurrence of a dog prediction
- Write that info into a separate field
- make that the dog prediction column categorial

Code

```
In [56]: dfc_image.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
tweet_id      2075 non-null int64
jpg_url       2075 non-null object
img_num       2075 non-null int64
p1            2075 non-null object
p1_conf       2075 non-null float64
p1_dog        2075 non-null bool
p2            2075 non-null object
p2_conf       2075 non-null float64
p2_dog        2075 non-null bool
p3            2075 non-null object
p3_conf       2075 non-null float64
p3_dog        2075 non-null bool
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 152.1+ KB

```

```
In [57]: #make sure the priority is correct in the columns
```

```
dfc_image.query('p1_conf < p2_conf | p2_conf < p3_conf | p1_conf < p3_conf')
```

```
Out[57]: Empty DataFrame
```

```

Columns: [tweet_id, jpg_url, img_num, p1, p1_conf, p1_dog, p2, p2_conf, p2_dog, p3, p3_
Index: []

```

```
In [58]: #merge thre dog columns into one according to priority
df_p1 = dfc_image[dfc_image.p1_dog == True].filter(items=['tweet_id', 'jpg_url', 'p1', 'p1_conf'])
df_p1 = df_p1.rename(columns={'p1': 'dog', 'p1_conf': 'conf'})
df_p2 = dfc_image.query('p1_dog == False & p2_dog == True').filter(items=['tweet_id', 'jpg_url', 'p2', 'p2_conf'])
df_p2 = df_p2.rename(columns={'p2': 'dog', 'p2_conf': 'conf'})
df_p3 = dfc_image.query('p1_dog == False & p2_dog == False & p3_dog == True').filter(items=['tweet_id', 'jpg_url', 'p3', 'p3_conf'])
df_p3 = df_p3.rename(columns={'p3': 'dog', 'p3_conf': 'conf'})
dfc_image = pd.concat([df_p1, df_p2, df_p3])
```

```
In [59]: dfc_image.dog = dfc_image.dog.astype('category')
```

Test

```
In [60]: dfc_image.head(5)
```

```
Out[60]:
```

	tweet_id	jpg_url	dog	conf
0	666020888022790149	https://pbs.twimg.com/media/CT4udnOWwAA0aMy.jpg	Welsh_springer_spaniel	0.465074
1	666029285002620928	https://pbs.twimg.com/media/CT42GRgUYAA5iDo.jpg	redbone	0.506826
2	666033412701032449	https://pbs.twimg.com/media/CT4521TWwAEvMyu.jpg	German_shepherd	0.596461
3	666044226329800704	https://pbs.twimg.com/media/CT5Dr8HUEAA-lEu.jpg	Rhodesian_ridgeback	0.408143
4	666049248165822465	https://pbs.twimg.com/media/CT5IQmsXIAAKY4A.jpg	miniature_pinscher	0.560311

```
In [61]: dfc_image.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1751 entries, 0 to 2026
Data columns (total 4 columns):
tweet_id    1751 non-null int64
jpg_url     1751 non-null object
dog         1751 non-null category
conf        1751 non-null float64
dtypes: category(1), float64(1), int64(1), object(1)
memory usage: 62.3+ KB
```

```
In [62]: df_image.query('tweet_id == 666044226329800704')
```

```
Out[62]:
```

	tweet_id	jpg_url
3	666044226329800704	https://pbs.twimg.com/media/CT5Dr8HUEAA-lEu.jpg

	img_num		p1	p1_conf	p1_dog	p2	p2_conf	p2_dog	\
3	1	Rhodesian_ridgeback	0.408143	True	redbone	0.360687	True		
			p3	p3_conf	p3_dog				
3		miniature_pinscher	0.222752	True					

Define join together dfc_twarchive x df_tweet add image info if it exists from dfc_image

Code

```
In [63]: df_tweet.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2345 entries, 0 to 2344
Data columns (total 3 columns):
tweet_id          2345 non-null int64
retweet_count     2345 non-null int64
favorite_count    2345 non-null int64
dtypes: int64(3)
memory usage: 55.0 KB
```

```
In [64]: j1 = pd.merge(dfc_twarchive, df_tweet, on='tweet_id')
         #filter out dogs where we don't know the type
         df_master = pd.merge(j1, dfc_image, on='tweet_id')
```

```
In [66]: df_master.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1664 entries, 0 to 1663
Data columns (total 14 columns):
tweet_id          1664 non-null int64
timestamp         1664 non-null object
source            1664 non-null category
text              1664 non-null object
expanded_urls     1664 non-null object
rating_numerator  1664 non-null int64
rating_denominator 1664 non-null int64
name              1203 non-null object
dog_stage         249 non-null category
retweet_count     1664 non-null int64
favorite_count    1664 non-null int64
jpg_url           1664 non-null object
dog               1664 non-null category
conf              1664 non-null float64
dtypes: category(3), float64(1), int64(5), object(5)
memory usage: 167.1+ KB
```

```
In [67]: df_master.sample(1)
```

```
Out[67]:
```

	tweet_id	timestamp	source		text		expanded_urls	rating_numerator		rating_denominator	name	dog_stage	retweet_count	favorite_count		jpg_url	dog	conf
1468	670782429121134593	2015-11-29 01:52:48 +0000	iPhone App		This dude slaps your girl's ass what do you do...		https://twitter.com/dog_rates/status/670782429...	5		10	None	NaN	847	1667		https://pbs.twimg.com/media/CU8ZDu9WwAADg3N.jpg	Chihuahua	0.952963

```
In [68]: df_master[df_master.jpg_url.isnull()].count()
```

```
Out[68]:
```

tweet_id	0
timestamp	0
source	0
text	0
expanded_urls	0
rating_numerator	0
rating_denominator	0
name	0
dog_stage	0
retweet_count	0
favorite_count	0
jpg_url	0
dog	0
conf	0
dtype:	int64

Test

```
In [70]: #prove that there is no duplicate jpg anymore.
```

```
df_master[df_master.duplicated(subset=['jpg_url'], keep='first')]
```

```
Out[70]: Empty DataFrame
```

```
Columns: [tweet_id, timestamp, source, text, expanded_urls, rating_numerator, rating_denominator, name, dog_stage, retweet_count, favorite_count, jpg_url, dog, conf]  
Index: []
```

Define set timezone to EST according to wikipedia that's where the twitter account owner studies at university

Code

```
In [74]: #convert object to datetime64
         #https://pandas.pydata.org/pandas-docs/stable/timeseries.html
         #localize to UTC and convert to EST
         #https://stackoverflow.com/questions/20689288/convert-pandas-columns-to-datetime64-ns
         #df_master['timestamp'] =
         df_master.timestamp = pd.to_datetime(df_master.timestamp).dt.tz_localize('UTC').dt.tz_convert('EST')
         #extract local hour
         df_master.create_HH24 = df_master.timestamp.dt.hour
```

TypeError Traceback (most recent call last)

```
<ipython-input-74-6c77b85dbeaf> in <module>()
    4 #https://stackoverflow.com/questions/20689288/convert-pandas-columns-to-datetime64-ns
    5 #df_master['timestamp'] =
----> 6 df_master.timestamp = pd.to_datetime(df_master.timestamp).dt.tz_localize('UTC').dt.tz_convert('EST')
    7 #extract local hour
    8 df_master.create_HH24 = df_master.timestamp.dt.hour
```

```
~\AppData\Local\Continuum\anaconda3\lib\site-packages\pandas\core\base.py in f(self, *args, **kwargs)
    209
    210         def f(self, *args, **kwargs):
--> 211             return self._delegate_method(name, *args, **kwargs)
    212
    213         f.__name__ = name
```

```
~\AppData\Local\Continuum\anaconda3\lib\site-packages\pandas\core\indexes\accessors.py in get_attr(self, name)
    131
    132         method = getattr(self.values, name)
--> 133         result = method(*args, **kwargs)
    134
    135         if not is_list_like(result):
```

```
~\AppData\Local\Continuum\anaconda3\lib\site-packages\pandas\util\_decorators.py in wrapper
    89         else:
    90             kwargs[new_arg_name] = new_arg_value
--> 91         return func(*args, **kwargs)
    92         return wrapper
    93         return _deprecate_kwarg
```

```
~\AppData\Local\Continuum\anaconda3\lib\site-packages\pandas\core\indexes\datetime.py in get_attr(self, name)
```

```

1835             new_dates = libts.tz_convert(self.asi8, 'UTC', self.tz)
1836         else:
-> 1837             raise TypeError("Already tz-aware, use tz_convert to convert.")
1838         else:
1839             tz = libts.maybe_get_tz(tz)

```

TypeError: Already tz-aware, use tz_convert to convert.

```
In [86]: df_master['create_HH24'] = df_master.timestamp.dt.hour.astype('category')
```

Test

```
In [88]: df_master.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1664 entries, 0 to 1663
Data columns (total 15 columns):
tweet_id          1664 non-null int64
timestamp         1664 non-null datetime64[ns, EST]
source            1664 non-null category
text              1664 non-null object
expanded_urls     1664 non-null object
rating_numerator  1664 non-null int64
rating_denominator 1664 non-null int64
name              1203 non-null object
dog_stage         249 non-null category
retweet_count     1664 non-null int64
favorite_count    1664 non-null int64
jpg_url           1664 non-null object
dog               1664 non-null category
conf              1664 non-null float64
create_HH24       1664 non-null category
dtypes: category(4), datetime64[ns, EST](1), float64(1), int64(5), object(4)
memory usage: 249.5+ KB

```

1.5 Storing, Analyzing, and Visualizing Data for this Project

Store the clean DataFrame(s) in a CSV file with the main one named **twitter_archive_master.csv**. If additional files exist because multiple tables are required for tidiness, name these files appropriately. Additionally, you may **store the cleaned data in a SQLite database** (which is to be submitted as well if you do).

Analyze and visualize your wrangled data in your wrangle_act.ipynb Jupyter Notebook. At least **three (3) insights and one (1) visualization** must be produced.

```
In [89]: df_master.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1664 entries, 0 to 1663
Data columns (total 15 columns):
tweet_id          1664 non-null int64
timestamp         1664 non-null datetime64[ns, EST]
source            1664 non-null category
text              1664 non-null object
expanded_urls     1664 non-null object
rating_numerator  1664 non-null int64
rating_denominator 1664 non-null int64
name              1203 non-null object
dog_stage         249 non-null category
retweet_count     1664 non-null int64
favorite_count    1664 non-null int64
jpg_url           1664 non-null object
dog               1664 non-null category
conf              1664 non-null float64
create_HH24       1664 non-null category
dtypes: category(4), datetime64[ns, EST](1), float64(1), int64(5), object(4)
memory usage: 249.5+ KB

```

```
In [91]: df_master.to_csv('twitter_archive_master.csv',index=False, encoding='utf-8')
```

```
In [5]: df_master = pd.read_csv('twitter_archive_master.csv')
```

```
In [6]: df_master.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1664 entries, 0 to 1663
Data columns (total 15 columns):
tweet_id          1664 non-null int64
timestamp         1664 non-null object
source            1664 non-null object
text              1664 non-null object
expanded_urls     1664 non-null object
rating_numerator  1664 non-null int64
rating_denominator 1664 non-null int64
name              1203 non-null object
dog_stage         249 non-null object
retweet_count     1664 non-null int64
favorite_count    1664 non-null int64
jpg_url           1664 non-null object
dog               1664 non-null object
conf              1664 non-null float64
create_HH24       1664 non-null int64
dtypes: float64(1), int64(6), object(8)
memory usage: 195.1+ KB

```

```
In [7]: df_master.source = df_master.source.astype('category')
df_master.dog_stage = df_master.dog_stage.astype('category')
df_master.dog = df_master.dog.astype('category')
df_master.create_HH24 = df_master.create_HH24.astype('category')
df_master.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1664 entries, 0 to 1663
Data columns (total 15 columns):
tweet_id          1664 non-null int64
timestamp         1664 non-null object
source            1664 non-null category
text              1664 non-null object
expanded_urls     1664 non-null object
rating_numerator  1664 non-null int64
rating_denominator 1664 non-null int64
name              1203 non-null object
dog_stage         249 non-null category
retweet_count     1664 non-null int64
favorite_count    1664 non-null int64
jpg_url           1664 non-null object
dog               1664 non-null category
conf              1664 non-null float64
create_HH24       1664 non-null category
dtypes: category(4), float64(1), int64(5), object(5)
memory usage: 156.5+ KB
```

```
In [8]: df_master.corr(method='pearson')
```

```
Out[8]:
```

	tweet_id	rating_numerator	rating_denominator	\
tweet_id	1.000000	0.550155		NaN
rating_numerator	0.550155	1.000000		NaN
rating_denominator	NaN	NaN		NaN
retweet_count	0.392921	0.317499		NaN
favorite_count	0.630534	0.420476		NaN
conf	0.103490	0.142090		NaN

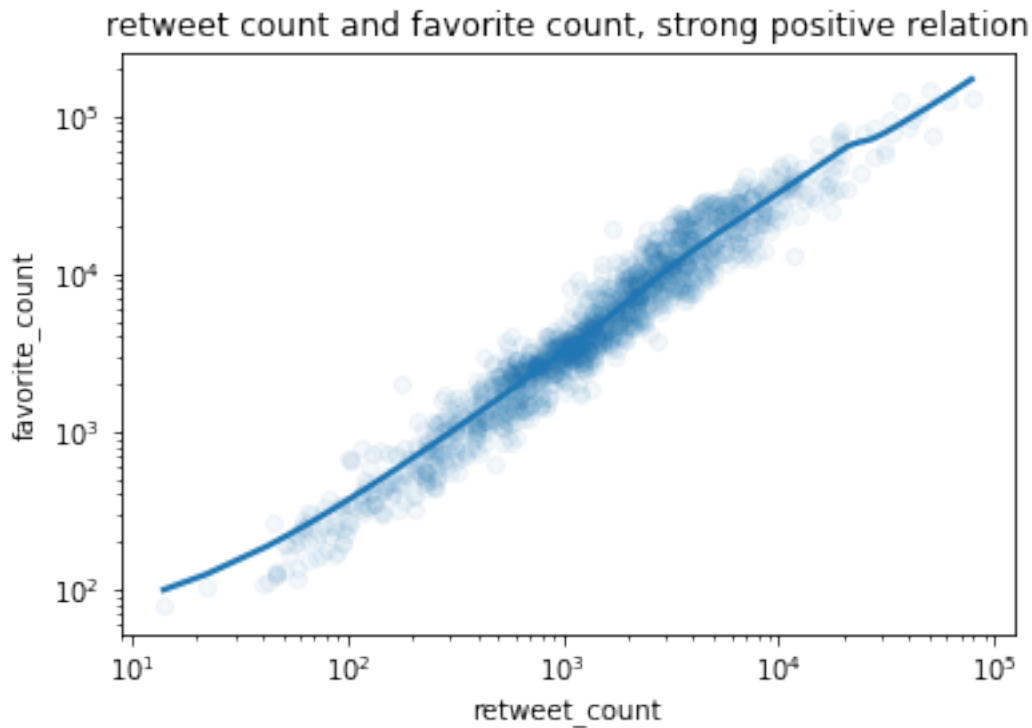
	retweet_count	favorite_count	conf
tweet_id	0.392921	0.630534	0.103490
rating_numerator	0.317499	0.420476	0.142090
rating_denominator	NaN	NaN	NaN
retweet_count	1.000000	0.917411	0.027693
favorite_count	0.917411	1.000000	0.059848
conf	0.027693	0.059848	1.000000

```
In [9]: # Initialize figure and ax
fig, ax = plt.subplots()
```

```

# Set the scale of the x-and y-axes
ax.set(xscale="log", yscale="log")
sns.regplot(x='retweet_count', y='favorite_count', data=df_master, ax=ax, scatter_kws={'
plt.title('retweet count and favorite count, strong positive relation')
plt.show()

```



```

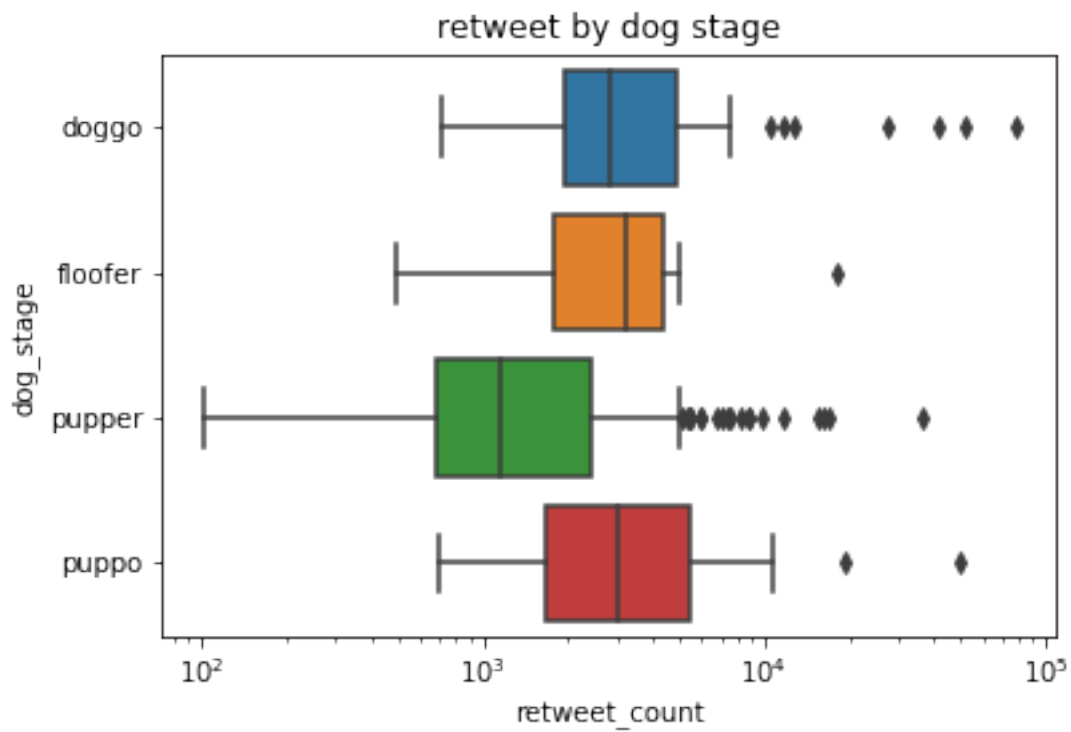
In [10]: # Create the boxplot
ax = sns.boxplot(x="retweet_count", y="dog_stage", data=df_master)
#ax = sns.boxplot(x="retweet_count", y="dog_stage", data=df_master)
#ax = sns.boxplot(x="rating_numerator", y="dog_stage", data=df_master)
# Initialize figure and ax
#fig, ax = plt.subplots()

# Set the `xlim`
ax.set(xscale="log")


# Set title
ax.set_title("retweet by dog stage")

# Show the plot
plt.show()

```



```
In [40]: #students work hard, ... after 10 in the morning.
ax = sns.countplot(x="create_HH24", data=df_master, palette="Greens_d")
ax.set_title('Busy twitter hours')
plt.show()
```



```
In [41]: pysqldf = lambda q: sqldf(q, globals())
```

```
q = """
SELECT
  dog
, count(*) as cnt
, avg(rating_numerator) as avg_rating_numerator
, avg(retweet_count) as avg_retweet_count
, avg(favorite_count) as avg_favorite_count
FROM df_master
GROUP BY 1
having cnt>10
ORDER BY avg_favorite_count desc
;
"""
```

```
In [42]: df_dog = pysqldf(q)
df_dog.head(10)
```

```
Out[42]:
```

	dog	cnt	avg_rating_numerator	avg_retweet_count	\
0	French_bulldog	31	11.193548	5179.774194	
1	whippet	11	10.181818	4995.909091	
2	standard_poodle	11	10.727273	5425.818182	
3	Samoyed	42	11.690476	4763.714286	